

Project 3: Collab and Compet

Learning Algorithm

1. Algorithm

I used a multi-agent DDPG to solve this Tennis environment. Each actor takes only its agent's state and outputs a vector representing an action chosen from a continuous action space. Critics, on the other hand, get states and actions of all agents and evaluate the action by computing a value function.

For the first 20,000 episodes, I trained the MADDPG agent using only random actions. After that, I add mean-zero Gaussian noise to actions and I reduce the scale of the noise over the course of training.

2. Hyperparameters

I used the following hyperparameters in my code:

- `BUFFER_SIZE = int(1e6)` # replay buffer size
- `GAMMA = 0.99` # discount factor
- `TAU = 1e-3` # for soft update of target
- `LR_ACTOR = 1e-4` # learning rate of the actor
- `LR_CRITIC = 1e-4` # learning rate of the critic
- `WEIGHT_DECAY = 1e-5` # L2 weight decay
- `UPDATE_EVERY = 20` # how often to update the network
- `BATCH_SIZE = 512` # minibatch size
- `NUM_OF_UPDATES = 1` # how many times to update

3. Model architectures

The architecture of Actor-Critic and the number of nodes in each layer is as follows:

Actor

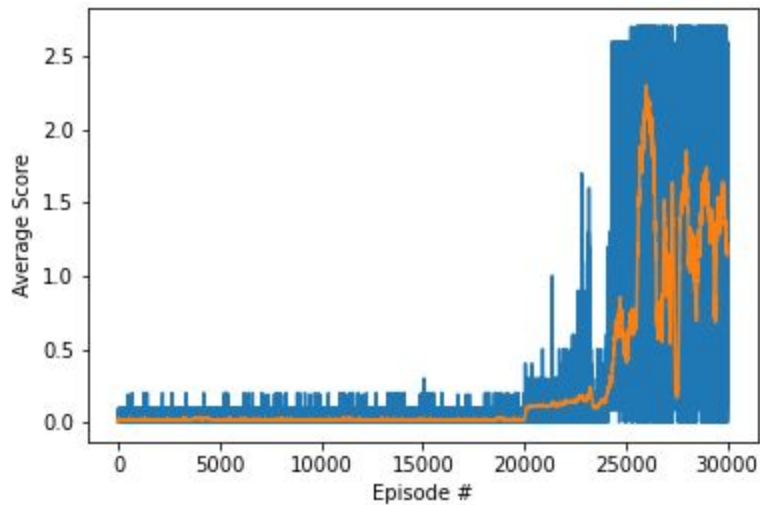
- Input layer = 24 (state_size)
- Hidden layer 1 (relu) = 128
- Hidden layer 2 (relu) = 128
- Output layer (tanh) = 4 (action_size)

Critic

- Input layer = 48 (state_size * num_agents)
- Hidden layer 1 (relu) = 128
- Concatenate = 132 (layer 1 size + (action size * num_agents))
- Hidden layer 2 (relu) = 128
- Output layer = 1

Plot of Rewards

The reward graph is shown below. My agent got an average score (over 100 episodes) of +0.5 in about 25000 episodes.



Ideas for Future work

1. [Trust Region Policy Optimization \(TRPO\)](#)
2. [Truncated Natural Policy Gradient \(TNPG\)](#)
3. [Proximal Policy Optimization \(PPO\)](#)
4. [Distributed Distributional Deterministic Policy Gradients \(D4PG\)](#)